Document Classification by Inversion of Distributed Language Representations

Matt Taddy

University of Chicago Booth School of Business

taddy@chicagobooth.edu

Abstract

There have been many recent advances in the structure and measurement of *distributed* language models: those that map from words to a vector-space that is rich in information about word choice and composition. This vector-space is the distributed language representation.

The goal of this note is to point out that any distributed representation can be turned into a classifier through inversion via Bayes rule. The approach is simple and modular, in that it will work with any language representation whose training can be formulated as optimizing a probability model. In our application to 2 million sentences from Yelp reviews, we also find that it performs as well as or better than complex purpose-built algorithms.

1 Introduction

Distributed, or vector-space, language representations $\mathcal{V}$ consist of a location, or embedding, for every vocabulary *word* in $\mathbb{R}^K$, where K is the dimension of the latent representation space. These locations are learned to optimize, perhaps approximately, an objective function defined on the original text such as a likelihood for word occurrences.

A popular example is the Word2Vec machinery of Mikolov et al. (2013b). This trains the distributed representation to be useful as an input layer for prediction of words from their neighbors in a Skip-gram likelihood. That is, to maximize

$$\sum_{j \neq t, \ j=t-b}^{t+b} \log \mathrm{p}_{\mathcal{V}}(w_{sj} \mid w_{st}) \qquad (1)$$

summed across all words w_{st} in all sentences $\mathbf{w}_s$, where b is the skip-gram window (truncated by the ends of the sentence) and $\mathrm{p}_{\mathcal{V}}(w_{sj}|w_{st})$ is a neural network classifier that takes vector representations for w_{st} and w_{sj} as input (see Section 2).

Distributed language representations have been studied since the early work on neural networks (Rumelhart et al., 1986) and have long been applied in natural language processing (Morin and Bengio, 2005). The models are generating much recent interest due to the large performance gains from the newer systems, including Word2Vec and the Glove model of Pennington et al. (2014), observed in, e.g., word prediction, word analogy identification, and named entity recognition.

Given the success of these new models, researchers have begun searching for ways to adapt the representations for use in document classification tasks such as sentiment prediction or author identification. One naive approach is to use aggregated word vectors across a document (e.g., a document's average word-vector location) as input to a standard classifier (e.g., logistic regression). However, a document is actually an *ordered* path of locations through $\mathbb{R}^K$, and simple averaging destroys much of the available information.

More sophisticated aggregation is proposed in Socher et al. (2011; 2013), where recursive neural networks are used to combine the word vectors through the estimated parse tree for each sentence. Alternatively, Le and Mikolov's Doc2Vec (2014) adds document labels to the conditioning set in (1) and has them influence the skip-gram likelihood through a latent input vector location in $\mathcal{V}$. In each case, the end product is a distributed representation for every sentence (or document for Doc2Vec) that can be used as input to a generic classifier.

1.1 Bayesian Inversion

These approaches all add considerable model and estimation complexity to the original underlying distributed representation. We are proposing a simple alternative that turns fitted distributed language representations into document classifiers

without any additional modeling or estimation.

A typical language model is trained to maximize the likelihoods of single words and their neighbors. For example, the skip-gram in (1) represents conditional probability for a word's context (surrounding words), while the alternative CBOW Word2Vec specification (Mikolov et al., 2013a) targets the conditional probability for each word given its context. Although these objectives do not correspond to a full document likelihood model, they can be interpreted as components in a *composite likelihood*[1] approximation.

Use $\mathbf{w} = [w_1 \ldots w_T]'$ to denote a sentence: an ordered vector of words. The skip-gram in (1) yields the pairwise composite log likelihood[2]

$$\log \mathrm{p}_{\mathcal{V}}(\mathbf{w}) = \sum_{j=1}^{T} \sum_{k=1}^{T} \mathbb{1}_{[1 \le |k-j| \le b]} \log \mathrm{p}_{\mathcal{V}}(w_k | w_j). \tag{2}$$

In another example, Jernite et al. (2015) show that CBOW Word2Vec corresponds to the pseudolikelihood for a Markov random field sentence model.

Finally, given a sentence likelihood as in (2), document $d = \{\mathbf{w}_1, \ldots \mathbf{w}_S\}$ has log likelihood

$$\log \mathrm{p}_{\mathcal{V}}(d) = \sum_s \log \mathrm{p}_{\mathcal{V}}(\mathbf{w}_s). \tag{3}$$

Now suppose that your training documents are grouped by class label, $y \in \{1 \ldots C\}$. We can train *separate* distributed language representations for each set of documents as partitioned by y; for example, fit Word2Vec independently on each sub-corpus $D_c = \{d_i : y_i = c\}$ and obtain the labeled distributed representation map $\mathcal{V}_c$. A new document d has probability $\mathrm{p}_{\mathcal{V}_c}(d)$ if we treat it as a member of class c, and Bayes rule implies

$$\mathrm{p}(y|d) = \frac{\mathrm{p}_{\mathcal{V}_y}(d)\pi_y}{\sum_c \mathrm{p}_{\mathcal{V}_c}(d)\pi_c} \tag{4}$$

where π_c is our prior probability on class label c.

Thus distributed language representations trained separately for each class label yield directly a document classification rule via (4). This approach has a number of attractive qualities.

Simplicity: The inversion strategy works for any model of language that can (or its training can) be interpreted as a probabilistic model. This makes for easy implementation in systems that are already engineered to fit such language representations, leading to faster deployment and lower development costs. The strategy is also interpretable: whatever intuition one has about the distributed language model can be applied directly to the inversion-based classification rule. Inversion adds a plausible model for reader understanding on top of any given language representation.

Scalability: when working with massive corpora it is often useful to split the data into blocks as part of distributed computing strategies. Our model of classification via inversion provides a convenient top-level partitioning of the data. An efficient system could fit separate by-class language representations, which will provide for document classification as in this article as well as class-specific answers for NLP tasks such as word prediction or analogy. When one wishes to treat a document as unlabeled, NLP tasks can be answered through ensemble aggregation of the class-specific answers.

Performance: We find that, in our examples, inversion of Word2Vec yields lower misclassification rates than both Doc2Vec-based classification and the multinomial inverse regression (MNIR) of Taddy (2013b). We did not anticipate such outright performance gain. Moreover, we expect that with calibration (i.e., through cross-validation) of the many various tuning parameters available when fitting both Word and Doc 2Vec the performance results will change. Indeed, we find that all methods are often outperformed by phrase-count logistic regression with rare-feature up-weighting and carefully chosen regularization. However, the out-of-the-box performance of Word2Vec inversion argues for its consideration as a simple default in document classification.

In the remainder, we outline classification through inversion of a specific Word2Vec model and illustrate the ideas in classification of Yelp reviews. The implementation requires only a small extension of the popular `gensim` python library (Rehurek and Sojka, 2010); the extended library as well as code to reproduce all of the results in this paper are available on `github`. In addition, the yelp data is publicly available as part of the corresponding data mining contest at `kaggle.com`. See `github.com/taddylab/deepir` for detail.

---

[1] Composite likelihoods are a common tool in analysis of spatial data and data on graphs. They were popularized in statistics by Besag's (1974; 1975) work on the pseudolikelihood $-$ $\mathrm{p}(\mathbf{w}) \approx \prod_j \mathrm{p}(w_j|\mathbf{w}_{-j})$ $-$ for analysis of Markov random fields. See Varin et al. (2011) for a detailed review.

[2] See Molenberghs and Verbeke (2006) for similar pairwise compositions in analysis of longitudinal data.

2 Implementation

Word2Vec trains $\mathcal{V}$ to maximize the skip-gram likelihood based on (1). We work with the Huffman softmax specification (Mikolov et al., 2013b), which includes a pre-processing step to encode each vocabulary word in its representation via a binary Huffman tree (see Figure 1).

Each individual probability is then

$$\mathrm{p}_{\mathcal{V}}(w|w_t) = \prod_{j=1}^{L(w)-1} \sigma\Big(\mathrm{ch}\left[\eta(w, j+1)\right] \mathbf{u}_{\eta(w,j)}^{\top} \mathbf{v}_{w_t}\Big) \tag{5}$$

where $\eta(w, i)$ is the i^{th} node in the Huffman tree path, of length $L(w)$, for word w; $\sigma(x) = 1/(1 + \exp[-x])$; and $\mathrm{ch}(\eta) \in \{-1, +1\}$ translates from whether η is a left or right child to +/- 1. Every word thus has both input and output vector coordinates, $\mathbf{v}_w$ and $[\mathbf{u}_{\eta(w,1)} \cdots \mathbf{u}_{\eta(w,L(w))}]$. Typically, only the input space $\mathbf{V} = [\mathbf{v}_{w_1} \cdots \mathbf{v}_{w_p}]$, for a p-word vocabulary, is reported as the language representation – these vectors are used as input for NLP tasks. However, the full representation $\mathcal{V}$ includes mapping from each word to both $\mathbf{V}$ and $\mathbf{U}$.

We apply the `gensim` python implementation of Word2Vec, which fits the model via stochastic gradient descent (SGD), under default specification. This includes a vector space of dimension $K = 100$ and a skip-gram window of size $b = 5$.

2.1 Word2Vec Inversion

Given Word2Vec trained on each of C class-specific corpora $D_1 \ldots D_C$, leading to C distinct language representations $\mathcal{V}_1 \ldots \mathcal{V}_C$, classification for new documents is straightforward. Consider


Figure 1: Binary Huffman encoding of a 4 word vocabulary, based upon 18 total utterances. At each step proceeding from left to right the two nodes with lowest count are combined into a parent node. Binary encodings are read back off of the splits moving from right to left.

the S-sentence document d: each sentence $\mathbf{w}_s$ is given a probability under each representation $\mathcal{V}_c$ by applying the calculations in (1) and (5). This leads to the $S \times C$ matrix of sentence probabilities, $\mathrm{p}_{\mathcal{V}_c}(\mathbf{w}_s)$, and document probabilities are obtained

$$\mathrm{p}_{\mathcal{V}_c}(d) = \frac{1}{S} \sum_s \mathrm{p}_{\mathcal{V}_c}(\mathbf{w}_s). \tag{6}$$

Finally, class probabilities are calculated via Bayes rule as in (4). We use priors $\pi_c = 1/C$, so that classification proceeds by assigning the class

$$\hat{y} = \mathrm{argmax}_c \ \mathrm{p}_{\mathcal{V}_c}(d). \tag{7}$$

3 Illustration

We consider a corpus of reviews provided by Yelp for a contest on `kaggle.com`. The text is tokenized simply by converting to lowercase before splitting on punctuation and white-space. The training data are 230,000 reviews containing more than 2 million sentences. Each review is marked by a number of *stars*, from 1 to 5, and we fit separate Word2Vec representations $\mathcal{V}_1 \ldots \mathcal{V}_5$ for the documents at each star rating. The validation data consist of 23,000 reviews, and we apply the inversion technique of Section 2 to score each validation document d with class probabilities $\mathbf{q} = [q_1 \cdots q_5]$, where $q_c = \mathrm{p}(c|d)$.

The probabilities will be used in three different classification tasks; for reviews as

a. negative at 1-2 stars, or positive at 3-5 stars;

b. negative 1-2, neutral 3, or positive 4-5 stars;

c. corresponding to each of 1 to 5 stars.

In each case, classification proceeds by summing across the relevant sub-class probabilities. For example, in task *a*, $\mathrm{p}(\texttt{positive}) = q_3 + q_4 + q_5$. Note that the same five fitted Word2Vec representations are used for each task.

We consider a set of related comparator techniques. In each case, some document representation (e.g., phrase counts or Doc2Vec vectors) is used as input to logistic regression prediction of the associated review rating. The logistic regressions are fit under L_1 regularization with the penalties weighted by feature standard deviation (which, e.g., up-weights rare phrases) and selected according to the corrected AICc criteria (Flynn et al., 2013) via the `gamlr` R package of Taddy (2014). For multi-class tasks *b-c*, we use distributed Multinomial regression (DMR;

Figure 2: Out-of-Sample fitted probabilities of a review being *positive* (having greater than 2 stars) as a function of the true number of review stars. Box widths are proportional to number of observations in each class; roughly 10% of reviews have each of 1-3 stars, while 30% have 4 stars and 40% have 5 stars.

Taddy 2015) via the `distrom` R package. DMR fits multinomial logistic regression in a factorized representation wherein one estimates independent Poisson linear models for each response category. Document representations and logistic regressions are always trained using only the training corpus.

Doc2Vec is also fit via `gensim`, using the same latent space specification as for Word2Vec: $K = 100$ and $b = 5$. As recommended in the documentation, we apply repeated SGD over 20 re-orderings of each corpus (for comparability, this was also done when fitting Word2Vec). Le and Mikolov provide two alternative Doc2Vec specifications: distributed memory (DM) and distributed bag-of-words (DBOW). We fit both. Vector representations for validation documents are trained without updating the word-vector elements, leading to 100 dimensional vectors for each document for each of DM and DCBOW. We input each, as well as the combined 200 dimensional DM+DBOW representation, to logistic regression.

Phrase regression applies logistic regression of response classes directly onto counts for short 1-2 word 'phrases'. The phrases are obtained using `gensim`'s phrase builder, which simply combines highly probable pairings; e.g., `first_date` and `chicken_wing` are two pairings in this corpus.

MNIR, the multinomial inverse regression of Taddy (2013a; 2013b; 2015) is applied as implemented in the `textir` package for R. MNIR maps from text to the class-space of interest through a multinomial logistic regression of phrase counts onto variables relevant to the class-space. We apply MNIR to the same set of 1-2 word phrases used in phrase regression. Here, we regress phrase counts onto stars expressed numerically and as a 5-dimensional indicator vector, lead-

	a (NP)	b (NNP)	c (1-5)
W2V inversion	.099	**.189**	.435
Phrase regression	**.084**	.200	**.410**
D2V DBOW	.144	.282	.496
D2V DM	.179	.306	.549
D2V combined	.148	. 284	.500
MNIR	.095	.254	.480
W2V aggregation	.118	.248	.461

Table 1: Out-of-sample misclassification rates.

ing to a 6-feature multinomial logistic regression. The MNIR procedure then uses the $6 \times p$ matrix of feature-phrase regression coefficients to map from phrase-count to feature space, resulting in 6 dimensional 'sufficient reduction' statistics for each document. These are input to logistic regression.

Word2Vec aggregation averages fitted word representations for a single Word2Vec trained on all sentences to obtain a fixed-length feature vector for each review ($K = 100$, as for inversion). This vector is then input to logistic regression.

3.1 Results

Misclassification rates for each task on the validation set are reported in Table 1. Simple phrase-count regression is consistently the strongest performer, bested only by Word2Vec inversion on task b. This is partially due to the relative strengths of discriminative (e.g., logistic regression) vs generative (e.g., all others here) classifiers: given a large amount of training text, asymptotic efficiency of logistic regression will start to work in its favor over the finite sample advantages of a generative classifier (Ng and Jordan, 2002; Taddy, 2013c). However, the comparison is also unfair to Word2Vec and Doc2Vec: both phrase regres-

sion and MNIR are optimized exactly under AICc selected penalty, while Word and Doc 2Vec have only been approximately optimized under a single specification. The distributed representations should improve with some careful engineering.

Word2Vec inversion outperforms the other document representation-based alternatives (except, by a narrow margin, MNIR in task a). Doc2Vec under DBOW specification and MNIR both do worse, but not by a large margin. In contrast to Le and Mikolov, we find here that the Doc2Vec DM model does much worse than DBOW. Regression onto simple within- document aggregations of Word2Vec perform slightly better than any Doc2Vec option (but not as well as the Word2Vec inversion). This again contrasts the results of Le and Mikolov and we suspect that the more complex Doc2Vec model would benefit from a careful tuning of the SGD optimization routine.[3]

Looking at the fitted probabilities in detail we see that Word2Vec inversion provides a more useful document *ranking* than any comparator (including phrase regression). For example, Figure 2 shows the probabilities of a review being 'positive' in task a as a function of the true star rating for each validation review. Although phrase regression does slightly better in terms of misclassification rate, it does so at the cost of classifying many terrible (1 star) reviews as positive. This occurs because 1-2 star reviews are more rare than 3-5 star reviews and because words of emphasis (e.g. `very`, `completely`, and `!!!`) are used both in very bad and in very good reviews. Word2Vec inversion is the *only* method that yields positive-document probabilities that are clearly increasing in distribution with the true star rating. It is not difficult to envision a misclassification cost structure that favors such nicely ordered probabilities.

4 Discussion

The goal of this note is to point out inversion as an option for turning distributed language representations into classification rules. We are not arguing

for the supremacy of Word2Vec inversion in particular, and the approach should work well with alternative representations (e.g., Glove). Moreover, we are not even arguing that it will always outperform purpose-built classification tools. However, it is a simple, scalable, interpretable, and effective option for classification whenever you are working with such distributed representations.

References

Julian Besag. 1974. Spatial interaction and the statistical analysis of lattice systems. *Journal of the Royal Statistical Society, Series B.*

Julian Besag. 1975. Statistical analysis of non-lattice data. *The Statistician*, pages 179–195.

Cheryl Flynn, Clifford Hurvich, and Jefferey Simonoff. 2013. Efficiency for Regularization Parameter Selection in Penalized Likelihood Estimation of Misspecified Models. *Journal of the American Statistical Association*, 108:1031–1043.

Yacine Jernite, Alexander Rush, and David Sontag. 2015. A fast variational approach for learning Markov random field language models. In *Proceedings of the 32nd International Conference on Machine Learning (ICML 2015).*

Quoc V. Le and Tomas Mikolov. 2014. Distributed representations of sentences and documents. In *Proceedings of the 31 st International Conference on Machine Learning.*

Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. 2013a. Efficient estimation of word representations in vector space. *arXiv preprint arXiv:1301.3781.*

Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg S. Corrado, and Jeff Dean. 2013b. Distributed representations of words and phrases and their compositionality. In *Advances in Neural Information Processing Systems*, pages 3111–3119.

Geert Molenberghs and Geert Verbeke. 2006. *Models for discrete longitudinal data.* Springer Science & Business Media.

Frederic Morin and Yoshua Bengio. 2005. Hierarchical probabilistic neural network language model. In *Proceedings of the International Workshop on Artificial Intelligence and Statistics*, pages 246–252.

Andrew Y. Ng and Michael I. Jordan. 2002. On Discriminative vs Generative Classifiers: A Comparison of Logistic Regression and naive Bayes. In *Advances in Neural Information Processing Systems (NIPS).*

---

[3]Note also that the unsupervised document representations – Doc2Vec or the single Word2Vec used in Word2Vec aggregation – could be trained on larger unlabeled corpora. A similar option is available for Word2Vec inversion: one could take a single Word2Vec model trained on a large unlabeled corpora as a shared baseline (prior) and update separate models with additional training on each labeled sub-corpora. The representations will all be shrunk towards a baseline language model, but will differ according to distinctions between the language in each labeled sub-corpora.

Jeffrey Pennington, Richard Socher, and Christopher D. Manning. 2014. Glove: Global vectors for word representation. *Proceedings of the Empirical Methods in Natural Language Processing (EMNLP 2014)*, 12.

Radim Rehurek and Petr Sojka. 2010. Software Framework for Topic Modelling with Large Corpora. In *Proceedings of the LREC 2010 Workshop on New Challenges for NLP Frameworks*, pages 45–50.

David Rumelhart, Geoffrey Hinton, and Ronald Williams. 1986. Learning representations by back-propagating errors. *Nature*, 323:533–536.

Richard Socher, Cliff C. Lin, Chris Manning, and Andrew Y. Ng. 2011. Parsing natural scenes and natural language with recursive neural networks. In *Proceedings of the 28th international conference on machine learning (ICML-11)*, pages 129–136.

Richard Socher, Alex Perelygin, Jean Y. Wu, Jason Chuang, Christopher D. Manning, Andrew Y. Ng, and Christopher Potts. 2013. Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of the conference on empirical methods in natural language processing (EMNLP)*, volume 1631, page 1642.

Matt Taddy. 2013a. Measuring Political Sentiment on Twitter: Factor Optimal Design for Multinomial Inverse Regression. *Technometrics*, 55(4):415–425, November.

Matt Taddy. 2013b. Multinomial Inverse Regression for Text Analysis. *Journal of the American Statistical Association*, 108:755–770.

Matt Taddy. 2013c. Rejoinder: Efficiency and structure in MNIR. *Journal of the American Statistical Association*, 108:772–774.

Matt Taddy. 2014. One-step estimator paths for concave regularization. arXiv:1308.5623.

Matt Taddy. 2015. Distributed Multinomial Regression. *Annals of Applied Statistics*, To appear.

Cristiano Varin, Nancy Reid, and David Firth. 2011. An overview of composite likelihood methods. *Statistica Sinica*, 21(1):5–42.